A speed-accuracy tradeoff in children's processing of scalar implicatures

Rose M. Schneider
rschneid@stanford.edu
Department of Psychology
Stanford University

Michael C. Frank
mcfrank@stanford.edu
Department of Psychology
Stanford University

Abstract

Children's trouble with scalar implicatures – inferences from a weaker lexicalized description that a stronger alternative is true – is a puzzle in pragmatic development. Here, we explore reaction time as a measure of processing for scalar implicatures and reasoning about salient alternatives. In our analyses, we explore overall performance and reaction time patterns across development, finding evidence of a speed-accuracy tradeoff for the quantifiers "some" and "none." Motivated by these findings, we use a Drift Diffusion Model to explore the relationship between accuracy and reaction time in processing both scalar implicatures, and the quantifiers "some" and "none" more broadly. Overall, we find evidence that while children's performance in scalar implicature tasks seems to requires additional processing when reasoning about the quantifier scale.

Keywords: Pragmatics; development; language.

Introduction

As listeners attempting to comprehend language, we have available to us not only an utterance, but also the knowledge of what the speaker *could* have said. In fact, we frequently go beyond the literal sense of utterances, and use our knowledge about these alternatives to infer a speaker's intended meaning. In the case of *pragmatic implicatures* (Grice, 1975), a speaker employs a weaker literal description to imply that a stronger alternative is true. Thus, an adult listener would strongly infer from the statement "I enjoyed *some* of my winter break" that some (but not *all*) of my break was pleasant. This *scalar implicature* (SI) relies heavily on a knowledge of the relevant lexical alternatives in the quantifier scale <*none – some – all*>, as a listener must be able to contrast "some" with the stronger descriptor "all" to compute the implicature. While scalar implicatures are easily comprehended by adults, they pose a pragmatic challenge to children until surprisingly late in development (Horowitz & Frank, 2015; Katsos & Bishop, 2011; Papafragou & Musolino, 2003). What is the source of children's difficulties with scalar implicatures?

One possibile cause of children's scalar implicature failures is their knowledge of the relevant lexical alternatives, as proposed by the *Alternatives Hypothesis* (D. Barner & Bachrach, 2010; David Barner, Brooks, & Bale, 2011). This hypothesis predicts that if children do not have access to "some," they are unable to directly compare it to "all" in making the scalar implicature *some, but not all.* Due to varying measures and methods in previous scalar implicature research, empirically testing the Alternatives Hypothesis was quite difficult. Children displayed varying performance across these tasks depending on paradigm, syntactic construction of the implicature prompts, access to visual and lexical alternatives, age, and supportiveness of the task (Guasti et al., 2005; Horowitz & Frank, 2015; Noveck, 2001; Papafragou & Musolino, 2003; Papafragou & Tantalou, 2004).

In an attempt to reconcile these various accounts and test the Alternatives Hypothesis, Horowitz and Frank (2015) designed a simple referent selection paradigm that could be used across a broad age range (3–5 years) to explore lexicalized (scalar) implicatures. In this task, children saw three book covers, each featuring four familiar objects (Figure 1), and the experimenter described a book using a scalar (quantifier) description (e.g., "On the cover of my book, *none/some/all* of the pictures are cats."). Importantly, children had access to lexical alternatives (over the course of the study), as well as visual alternatives (within each trial).

With this supportive paradigm, Horowitz and Frank found that children struggled with the quantifiers "some" and "none" in relation to "all". Intriguingly, they found that performance between these two quantifiers was strongly bimodal and correlated: Children who failed on trials with the quantifier "some" similarly struggled with "none," and vice versa. While Horowitz & Frank's (2015) findings did provide some support for the Alternatives Hypothesis (with children's scalar implicature performance supported by access to alternatives), they observed that children were not able to capitalize on this knowledge to improve over the course of the study. Therefore, they hypothesized that there was another possible cause for children's poor scalar implicature computation.

Horowitz, Schneider, & Frank (in prep.) presented two hypotheses for children's observed patterns of performance, namely, lack of quantifier knowledge and developing inhibitory control (Horowitz et al., in prep.). There, we reasoned that the Alternatives Hypothesis necessitates familiarity with and ability to contrast alternatives on the quantifier scale <*none – some – all*>; if children's quantifier knowledge is absent or unestablished, it might lead to failures in making an implicature. Another possible cause of children's struggles with scalar implicatures might be that they have

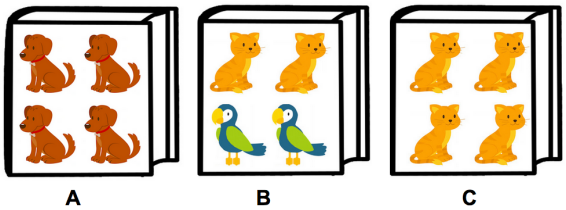

Figure 1: Example trial stimuli used in Horowitz and Frank (2015).

complete quantifier knowledge, but are unable to inhibit an impulse to choose a more salient alternative (e.g., choosing the book with the most cats upon hearing "On the cover of my book, *some* of the pictures are cats.")

We explored these two hypotheses in an individual differences task, running our scalar implicature (Horowitz & Frank, 2015) task in conjunction with quantifier-knowledge (D. Barner, Chow, & Yang, 2009) and inhibitory control (Zelazo, 2006) tasks. Overall, we found that children's scalar implicature performance was strongly correlated with quantifier knowledge, even when controlling for age. While inhibitory control was strongly correlated with age, it was not related to children's performance on the scalar implicature task (Horowitz et al., in prep.).

Although we found a significant relationship between quantifier knowledge and scalar implicature comprehension, it is likely that children's difficulties in this task has more than one source. Importantly, in the course of conducting this individual differences study, we observed that children who succeeded both in making a scalar implicature and comprehending "none" displayed increased response latencies. This increase in reaction time may be an important indicator of how children are using and evaluting the salient alternatives.

It is possible that children who have a fully-established quantifier scale must take additional time in using this information to process a scalar implicature. This hypothesis has some support in previous literature; Huang & Snedeker (2009) explored online measures scalar implicature processing through eye-tracking, and found a delay in children's location of the referent of a scalar implicature. The exact relationship between reaction time and accuracy using a supportive behavioral paradigm, however, is largely unknown. Here, we explore children's behavioral response latencies in an iPad adaptation of Horowitz and Frank's (2015) scalar implicature task. In this study, we explore our hypothesis that computing a scalar implicature might incur additional processing time for children as they contrast the relevant lexical alternatives to make a correct decision.

In our analyses, we explore overall accuracy and patterns of performance, as in (Horowitz & Frank, 2015; Horowitz et al., in prep.), and find that children not only struggle in making a scalar implicature, but also grapple with the quantifier "none" until fairly late in development. In examining reaction time patterns across all quantifier types, we find evidence of a speed-accuracy tradeoff associated with these two quantifiers, even later in development. Finally, we use a Drift Diffusion Model to our data to explore the source of this increased reaction time. Overall, our findings indicate that while quantifier knowledge is a key factor in successfully computing scalar implicatures, using this information to successfully compute a scalar implicature is particularly difficult, and seems to requires additional processing.

Method

In this study, we adapted a scalar mplicature paradigm developed by Horowitz and Frank (2015) for the iPad.[1] In addition to capturing detailed reaction time data, this version included more trials, and standardized prosody across all trials, in addition to a completely randomized design.[2]

Participants

Age group	N	Mean	Median	SD
3–3.5 years	24	3.27	3.27	0.14
3.5–4 years	35	3.78	3.73	0.15
4–4.5 years	25	4.28	4.28	0.15
4.5–5 years	30	4.76	4.76	0.15
5–6.5 years	24	5.55	5.56	0.36

Table 1: Age demographic information for all participants.

Table 1 shows the breakdown of age information for all participants. Included in analyses are 138, 138, 138, 138, 138 children out of a planned sample of 120 participants. We ran 20 additional children, who were excluded from analysis for low English language exposure or $<$50% of trials completed. Included in our sample were 79 females and 59 males. Based on (Horowitz & Frank, 2015; Horowitz et al., in prep.), the initially planned sample size was 96 children from 3–5 years. After collecting data from 57 participants, however, we observed significantly lower performance on implicature trials across all age groups, indicating that the iPad adaptation of the scalar implicature task was slightly more challenging for all children, and included an older age group of 24 5–6.5-year-olds.

Stimuli

The general format of the task was identical to (Horowitz & Frank, 2015), with the exception of added items for additional trials. The study was programmed in HTML, CSS, and JavaScript, and displayed to children on a full-sized iPad. Each trial displayed three book covers, each containing a set of four familiar objects (Figure 1). Each trial allowed 2.5s for children to visually inspect the three book covers, before the experiment played the trial prompt (e.g., "On the cover of my book, *none* of the pictures are cats."). Each trial was randomized, with the exception that similar items were displayed together (e.g., food, clothing). Each session involved 30 trials, with 10 trials per quantifier-type ("all", "some", and "none"). In our randomization, quantifier triad order, items (within category), target item, and quantifier, were randomized for all participants.

[1]The full experiment can be viewed online at `https://rosemschneider.github.io/tablet_exp/si_tablet.html`.

[2]All of our data, processing, experimental stimuli, and analysis code can be viewed in the version control repository for this paper at: `https://github.com/rosemschneider/SI_tablet`.

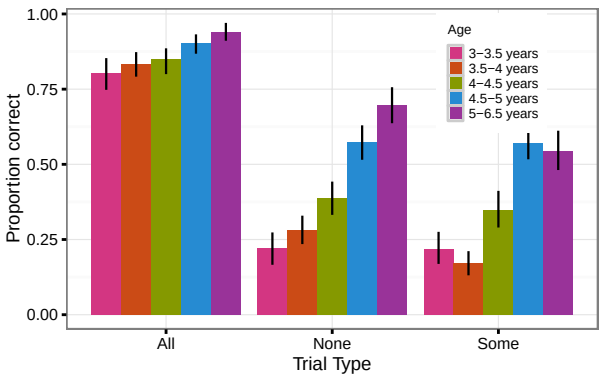

Figure 2: Children's overall accuracy for each quantifier type. Bars show mean performance for each age group. Error bars are 95 percent confidence intervals computed by non-parametric bootstrap.

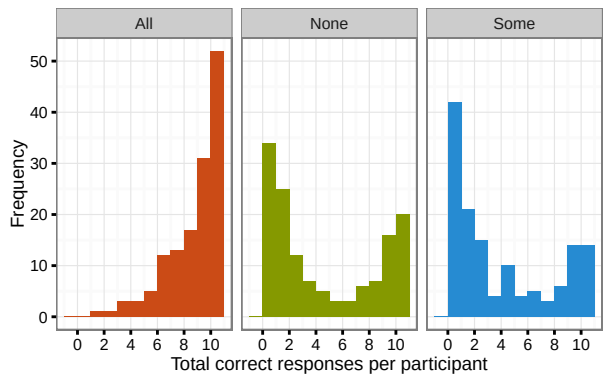

Figure 3: Frequency histogram of correct responses for each trial type, across all participants.

Procedure

Sessions took place individually in a small testing room away from either the museum or the classrooom. To familiarize children to the iPad, each session began with the "dot game," which required them to press dots on the screen as fast as possible.

After the dot game, the experimenter introduced them to "Hannah," a cartoon character who wanted to play a guessing game with her books. The experimenter explained that Hannah would show the child three books, and would give *one* hint about which book she had in mind. The experimenter emphasized that Hannah would only give one hint, so they had to listen carefully. Children then saw a practice trial with three books featuring a refrigerator, a TV, and a couch. After 2.5s, a female voice said "On the cover of my book there's a TV." Once children correctly made their selection, a green box appeared around the selection. Children moved trials along at their own pace by pressing a green button that appeared after they had made their selection.

Reaction times were measured from the onset of the target word. Each audio clip used the same three frames (e.g., "On the cover of my book, *some* of the pictures...") so that prodosdy was emphasized equally across all trials. The average length of each audio clip (including target item phrase, e.g., "...are cats") was approximately 6s. In all, there were 270 different target items and audio clips. Children could only make one selection. If a child was not paying attention, or if she did not hear Hannah's prompt, the experimenter repeated it, matching the original prosody.

Results

In analyzing the results, we excluded any trials in which reaction time exceeded fifteen seconds, which indicated that the child had missed the prompt, or was not paying attention. After this initial cut, we excluded responses outside three standard deviations of the log of mean reaction time. This cleaning process resulted in a data loss of 85 trials (2.14%). Our planned analyses for this study included explorations of accuracy, reaction time, and fitting a drift diffusion model (Ratcliff & Rouder, 1998) to our reaction time data.

Accuracy

Overall accuracy In our first planned analysis, we explored children's overall patterns of accuracy for each quantifier type. Figure 2 shows children's for each trial type, split by age group. For each age group, we saw significantly lower accuracy for the quantifiers "some" and "none" in comparison to "all" in independent t-tests within each age group ($p <$ .01 for all tests). These results replicate our previous findings using this paradigm (Horowitz & Frank, 2015; Horowitz et al., in prep.), indicating that our iPad version was a successful and appropriate adaptation of the scalar implicature task.

One difference from the previous results was in implicature trials. We found that children aged 3–5 years performed significantly lower on "some" (implicature) trials in this task in comparison with (Horowitz et al., in prep.) in independent t-tests ($p <$ .009 for all tests). This difference is most likely due to the fact that our adaptation relied strictly on verbal communication, rather than other social and nonverbal cues.

Statistical modeling In exploring children's signficantly lower performance on "some" and "none" trials, we ran a logistic mixed effects model predicting correct response as an interaction of age and trial type, with random effects of trial type and participant.[3] We found that performance was significantly lower on "some" (β = -6.98, $p <$ .0001) and "none" trials (β = -9.55, $p <$ .0001). We also found a signficant interaction between age and trial type on "none" trials(β = 1.53, $p <$ .0001), indicating that children's performance with this difficult quantifier increased with age.

Correlation between "some" and "none" In previous research, a strong correlation has been found on children's performance with the quantifiers "some" and "none" (Horowitz

[3]Mixed effects model fit in R using the lme4 package. The model specifications were as follows: `correct  age * trial type + (trial type | subject id)`.

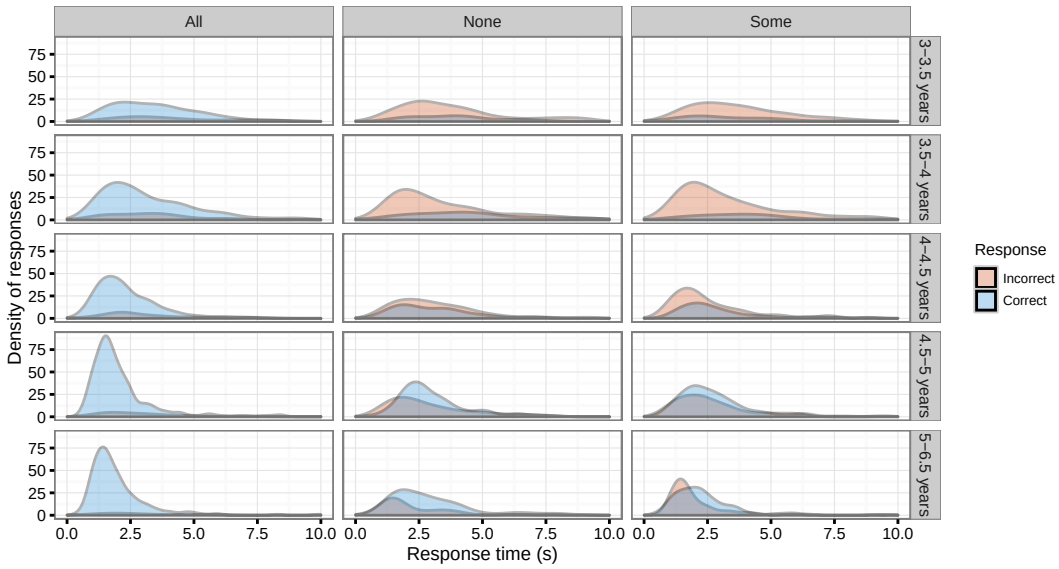

Figure 4: Density plots of reaction times for correct and incorrect responses on each trial type, split by age.

& Frank, 2015; Horowitz et al., in prep.). Figure 3 shows distributions of correct responses for all trial types. Once again, we found correlated performance between these two quantifiers (r = 0.49, $p < .0001$). In running Hartigan's diptest for bimodality on these two quantifiers, we found significant bimodal distributions for "some" (D = 0.08, $p < .0001$) and "none" trials (D = 0.11, $p < .0001$). While the diptest also signficantly rejected unimodality for "all" trials, this is most likely due to the distribution's long tail. The observed significantly bimodal performance on "some" and "none" trials, however, indicate that while children's accuracy is significantly worse for these quantifiers, they make responses in meaningful manner over the course of the study. In fact, although children's performance in general was quite low on these trials, there were a number of children who got the majority of these trials correct ("Some": N = 28; "None": N = 36).

Reaction time analyses

Response latencies are a largely-unexplored aspect of children's ability to compute scalar implicatures. We hypothesized that children's reaction times on this task may reflect processing of a scalar prompt, and weighing quantifier alternatives. It is possible that this measure may provide a clue as to the nature of children's correlated struggles with these terms. In recording reaction times, we began recording from the onset of the target nouns. Here, we explore overall trends in reaction times across this task, and the relationship between response latencies and accuracy on this task.

Developmental reaction time distribution Figure 4 shows the density of reaction times for each quantifier, faceted by age group and trial type. Overall, we found that reaction time was negatively correlated with age (r = -0.25, $p < .0001$). In exploring the relationship between accuracy and reaction time, we found preliminary evidence of a speed-accuracy tradeoff across all trial types. Figure 4 also shows evidence for increased response latencies associated with correct responses for the quantifiers "some" and "none". Overall, we found that children were largely consistent in their performance across the course of the study in reaction time correlations between reaction times for "all" and "some" (r = 0.76, $p < .0001$), "all" and "none" (r = 0.65, $p < .0001$), and "some" and "none" trials (r = 0.7, $p < .0001$).

Statistical modeling We next turned to the relationship betwen age, reaction time, and accuracy. Our initial hypothesis was that successfully computing quantifiers might require additional processing time to contrast salient alternatives. In exploring this, we ran a planned linear mixed effects model predicting the log of reaction time as an interaction of age, trial number, and trial type with a random effect of trial type.[4] We found a main effect of trial number, with reaction times decreasing over the course of the study ((β = -0.1, $p < .00001$)), but found longer reaction times on "none" (β = 0.22, $p < .00001$), and "some" trials (β = 0.1, $p < .00001$). Interestingly, we also found an interaction between age and trial type, such that reaction times on "none" trials increased with age relative to "all" trials (β = -0.01, $p < .02$), and marginally (but not significantly) increased on "some" trials (β = 0.14, p = .38).

This interaction is particularly intriguing because in our previous accuracy model, we found increased performance on these trial types. While we find that older children are tak-

[4]Mixed effects model fit in R using the lme4 package. The model specifications were as follows: `log(reaction time) scale(age) * log(trial number) + scale(age) * trial type + (trial type | subject id)`. We calculated p values by treating the t statistic as if it were a z statistic Barr, Levy, Scheepers, & Tily (2013).

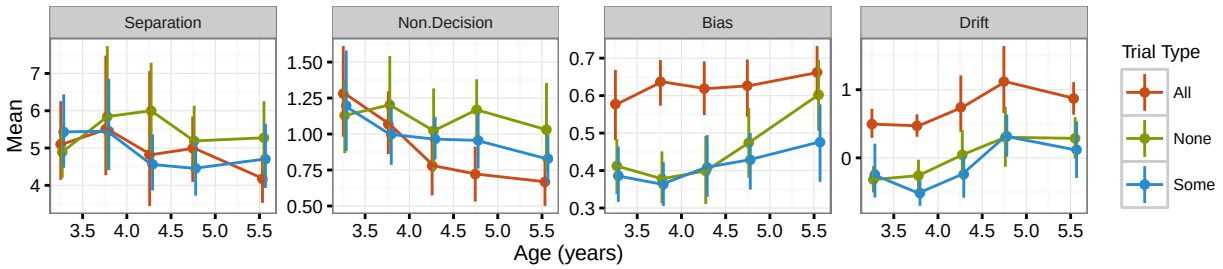

Figure 5: Parameter estimates for drift diffusion model, split by age and trial type. Error bars are 95 percent confidence intervals computed by nonparametric bootstrap.

ing longer to respond to these trial types, they are more likely to answer correctly, even though reaction time is negatively correlated with age. This seems to indicate that there is a speed-accuracy tradeoff for these quantifiers.

Drift diffusion models

Our preliminary analysis of children's reaction times in this task indicated greater response latencies associated with success on "some" and "none" trials. This suggests that children may be taking more time to actively compare and contrast scalar alternatives as they become more familiar with the quantifier scale. A drift diffusion model (DDM) can provide a more detailed view of the relationship between accuracy and reaction time in behavioral tasks (Milosavljevic, Malmaud, Huth, Koch, & Rangel, 2010; Ratcliff & Rouder, 1998).

Parameter estimation In DDM, a behavioral response (a correct or incorrect choice) is the result of noisy data accumulation through a diffusion process (operationalized by response time) (Ratcliff & Rouder, 1998). Responses have *separation boundaries* that are dependent on the amount of information needed to initate a response, and *drift rate* formalizes the rate of data accumulation (Ratcliff & Rouder, 1998). An additional parameter of DDM is *nondecision*, which is the amount of time between encoding the stimuli, and initiating response. Finally, different responses may have a *bias*, or different starting point in the diffusion process, dependent on the stimuli (Ratcliff & Rouder, 1998).

Although a DDM is traditionally used in two-alternative forced-choice tasks, here we are estimating the drift process between a correct and incorrect choice, with two options in each trial being "incorrect," and only one being consistent with the target noun and quantifier. In fitting a DDM to our data in this way, we estimated parameters for each subject across all three trial types ("all", "some", and "none") using the RWiener package. We then aggregated across subjects to obtain means and confidence intervals for each age group. Figure 5 shows the parameter estimates for each age group, split by trial type.

Using the DDM, we can explore in more depth the speed-accuracy tradeoff we observed in our statistical models. In our separation boundary estimates, we did not see any significant effect of trial type on boundary estimates, indicating that roughly the same amount of information needs to be accumulated to make a decision in each trial type. We also did not find an effect of trial type on drift rate, although in a mixed effects model we did find a main of effect of age on drift rate (β 0.24, p = .01).[5] We did not find that children accumulated data significantly faster for "some" and "none" trials.

In our nondecision parameter estimates, we again did not find a significant effect of quantifier type, although in a mixed effects model we found a main effect of age (β = -0.28, $p <$ .00001), and a significant interaction between age and "none" trials (β = 0.24, p = .01). Thus, the model predicts that although older children have lower nondecision times, these times are higher for older children on "none" trials (in comparison to "all") trials.

Finally, in our bias estimates, we found high bias across all age groups on "all" trials, and markedly lower estimates for "some" trials. Bias estimates for "none" trials increases over development, and we found a significant interaction between age and trial type on "none" trials (β = 0.07, p = .023). This higher bias for "none" versus "some" may be a result of the additional step that children must make in computing this quantifier. Additionally, previous research indicates that when making a scalar implicature, "none" is a particularly salient and compelling alternative (Franke, 2014).

Taken together, the parameter estimates from our DDM indicate that as children develop (and become more familiar with the quantifier scale), they are more likely to respond correctly in our scalar implicature task (especially on "none"). Younger children are more likely to respond incorrectly, mostly due to a low rate of data accumulation and a high separation boundary. Children's bias to respond correctly for all trial types seems to increase with age, and nondecision appears to decrease. While these findings do not fully explain the interaction between response time and accuracy we observed in our behavioral data, they do indicate some developmental differences in processing underying children's development in computing implicatures.

[5]The specifications for all parameter models are as follows: `Parameter Value  age * trial type + (1 | subject ID)`

General Discussion

Our primary question in this study centered on whether success in making scalar implicatures requires increased response latencies to make use of relevant scalar alternatives. We adapted a previously validated scalar implicature task (Horowitz & Frank, 2015; Horowitz et al., in prep.) for the iPad to explore the relationship between reaction time and accuracy.

In our analyses, we replicated previous patterns of performance in Horowitz & Frank (2015) and Horowitz et al. (in prep.). We found that children were overall less accurate when evaluting the quantifiers "some" and "none" in comparison to "all," but that their performance increased over development. We again found evidence of bimodal and correlated performance on these two quantifier types, suggesting a common source of difficulty. Additionally, we discovered in a statistical model that although children were more likely make an incorrect response on "some" and "none" trial, their performance on these trials signficantly increased with age.

In our extension of this paradigm, we collected reaction time data for these quantifier types to investigate the relationship between reaction time and accuracy. In our reaction time analyses, we found evidence of a speed-accuracy tradeoff, as well as an interaction between reaction time and age, with older children taking a slightly longer time to respond to these trials, but ultimately being more accurate. These findings motivated our decision to fit our data to a Drift Diffusion Model.

Using a DDM, we explored reaction time and accuracy patterns in more depth. Overall, the model predicted that young children accumulate data quickly, but are more likely to make an incorrect response on "some" and "none" trials. With age, children do not accumulate data more quickly, but are more biased to make a correct response on these trials. Thus, it appears that success in making a scalar implicature is not only the result of being familiar with the quantifier scale, but also taking additional time to process and integrate in a pragmatic framework in order to make a response.

Our work contributes to the existing literature in utilizing a novel method to collect accurate and detailed reaction time data on a scalar implicature task. Response latencies are an important indicator of the pragmatic challenges that children face in processing implicatures. Additionally, our findings replicate previous work, providing evidence for the appropriateness of this paradigm in targeting scalar implicatures. Further, our larger sample size, increased number of trials, and randomized design strengthen our analytical power, and allow for more detailed inferences from the data. Our work supports not only the hypthesis that children must be familiar with the quantifier scale in order to make an implicature, but also provides preliminary evidence that doing correctly doing so may require additional processing time.

Taken together, our work suggests that there is a meaningful relationship between children's accuracy and reaction times in making scalar implicatures throughout development. From our results, the relationship between children's quantifier knowledge, processing speed, and scalar implicature computation is unclear, and future work should test these links more explicitly. Our work suggests, however, that comparing a speaker's statement to possible alternatives in order to make an implicature is an active process for children, and does seem to result in a speed-accuracy tradeoff.

Acknowledgements

Special thanks to Bing Nursery School, the San Jose Children's Discovery Museum, Veronica Cristiano, Rachel Walker, and Tamara Mekler for their help with data collection, and to Kara Weisman and Ann Nordmeyer for their assistance creating stimuli.

References

Barner, D., & Bachrach, A. (2010). Inference and exact numerical representation in early language development. *Cognitive Psychology*, *60*(1), 40–62.

Barner, D., Brooks, N., & Bale, A. (2011). Accessing the unsaid: The role of scalar alternatives in childrens pragmatic inference. *Cognition*, *118*(1), 84–93.

Barner, D., Chow, K., & Yang, S. (2009). Finding one's meaning: A test of the relation between quantifiers and integers in language development. *Cognitive Psychology*, *58*(2), 195–219.

Barr, D. J., Levy, R., Scheepers, C., & Tily, H. (2013). Random effects structure for confirmatory hypothesis testing: Keep it maximal. *Journal of Memory and Language*, *68*(3), 255–278.

Franke, M. (2014). Typical use of quantifiers: A probabilistic speaker model. In *Proceedings of the 36th annual conference of the cognitive science society* (pp. 487–492).

Grice, H. (1975). *Logic and conversation* (pp. 41–58).

Guasti, M., Chierchia, G., Crain, S., Foppolo, F., Gualmini, A., & Meroni, L. (2005). Why children and adults sometimes (but not always) compute implicatures. *Language and Cognitive Processes*, *20*(5), 667–696.

Horowitz, A., & Frank, M. C. (2015). Sources of developmental change in pragmatic inferences about scalar terms. In *Proceedings of the 37th annual conference of the cognitive science society.*

Horowitz, A., Schneider, R. M., & Frank, M. C. (in prep.). The trouble with quantifiers: Children's difficulties with "some" and "none".

Huang, Y. T., & Snedeker, J. (2009). Online interpretation of scalar quantifiers: Insight into the semantics-pragmatics interface. *Cognitive Psychology*, *58*(3), 376–415.

Katsos, N., & Bishop, D. (2011). Pragmatic tolerance: Implications for the acquisition of informativeness and implicature. *Cognition*, *120*(1), 67–81.

Milosavljevic, M., Malmaud, J., Huth, A., Koch, A., & Rangel, A. (2010). Drift diffusion model can account for accuracy and reaction time of value-based choices under high and low time pressure. *Judgment and Decision Making*, *5*(6), 437–449.

Noveck, I. (2001). When children are more logical than adults: Experimental investigations of scalar implicature. *Cognition*, *78*(2), 165–188.

Papafragou, A., & Musolino, J. (2003). Scalar implicatures: Experiments at the semantics-pragmatics interface. *Cognition*, *86*(3), 253–282.

Papafragou, A., & Tantalou, N. (2004). Children's computation of implicatures. *Language Acquisition*, *12*, 71–82.

Ratcliff, R., & Rouder, J. (1998). Modeling response times for two-choice decisions. *Psychologial Science*, *9*(5), 347–356.

Zelazo, P. D. (2006). The dimensional change card sort (dCCS): A method of assessing executive function in children. *Nature Protocols*, *1*, 297–301.