

A speed-accuracy tradeoff in children's processing of scalar implicatures

Rose M. Schneider

rschneid@stanford.edu

Department of Psychology

Stanford University

Michael C. Frank

mcf Frank@stanford.edu

Department of Psychology

Stanford University

Abstract

Children's trouble with scalar implicatures – inferences from a weaker lexicalized description that a stronger alternative is true – is a puzzle in pragmatic development. Previous research indicates that children's failures in processing scalar implicatures may be rooted in an unestablished quantifier scale, with children who struggle with the quantifier "some" being unable to contrast different quantifiers to make the implicature. However, the source of this failure is unclear. Here, we explore reaction time as a measure of processing for scalar implicatures and reasoning about salient alternatives. In our analyses, we explore overall performance and reaction time patterns across development, finding that increased reaction times and accuracy for the quantifiers "some" and "none." Motivated by these findings, we use a Drift Diffusion Model to explore the relationship between accuracy and reaction time in processing both scalar implicatures, and the quantifiers "some" and "none" more broadly. Overall, we find evidence that while children's performance in scalar implicature tasks seems to require additional processing when reasoning about the quantifier scale.

Keywords: Pragmatics; development; language.

Introduction

As listeners attempting to comprehend language, we have available to us not only verbalized linguistic information, but also the knowledge of what the speaker *could* have said. In fact, we frequently go beyond the literal sense of utterances, and use knowledge of these alternatives to infer a speaker's intended meaning. In the case of *pragmatic implicatures* (Grice, 1975), a speaker employs a weaker literal description to imply that a stronger alternative is true. Thus, an adult listener would strongly infer from the statement "I enjoyed *some* of my winter break" that some (but not *all*) of my break was pleasant. This *scalar implicature* (SI) relies heavily on a knowledge of the relevant lexical alternatives in the quantifier scale $\langle \textit{none} - \textit{some} - \textit{all} \rangle$, as a listener must be able to contrast these alternatives to compute the implicature. While scalar implicatures are easily comprehended by adults, they pose a pragmatic challenge to children until fairly late in development (Horowitz & Frank, 2015; Katsos & Bishop, 2011; Papafragou & Musolino, 2003). What is the source of children's difficulties with scalar implicatures?

One possible cause of children's scalar implicature failures is their knowledge of the relevant lexical alternatives, as proposed by the *Alternatives Hypothesis* (D. Barner & Bachrach, 2010; David Barner, Brooks, & Bale, 2011). For example, this hypothesis predicts that if children do not have access to "some," they are unable to directly compare it to "all" in making the scalar implicature *some*, but not *all*. Across research on the development of scalar implicature processing, however, children exhibit varying performance, de-

pending on the paradigm, syntactic construction of the implicature prompts, access to visual and lexical alternatives, and age (Guasti et al., 2005; Horowitz & Frank, 2015; Noveck, 2001; Papafragou & Musolino, 2003; Papafragou & Tantalou, 2004; Stiller, Goodman, & Frank, 2014). Due to these varying measures and methods, making comparisons across these datasets, and empirically testing the Alternatives Hypothesis was quite difficult.

In an attempt to reconcile these various accounts and test the Alternatives Hypothesis, Horowitz and Frank (2015) designed a simple referent selection paradigm that could be used across a broad age range (3–5 years) to explore lexicalized (scalar) implicatures. In this task, children saw three book covers, each featuring four familiar objects (Figure 1), and the experimenter described a book using scalar (quantifier) description (e.g., "On the cover of my book, *none/some/all* of the pictures are cats."). Importantly, children had access to lexical alternatives (over the course of the study), as well as visual alternatives (within each trial).

With this supportive paradigm, Horowitz and Frank found children struggled with the quantifiers "some" and "none" in relation to "all", although performance increased across development. Intriguingly, they found that performance between these two quantifiers was strongly bimodal and correlated: Children who failed on trials with the quantifier "some" similarly struggled with "none," and vice versa. While Horowitz & Frank's (2015) findings did provide some support for the Alternatives Hypothesis (with children's scalar implicature performance supported by access to lexical alternatives), they observed that children were not able to use these alternatives to improve over the course of the study. Therefore, they hypothesized that there was another possible cause for children's poor scalar implicature computation.

These early failures in computing scalar implicatures, and the related performance processing the quantifiers "some" and "none" observed in (Horowitz & Frank, 2015) presented two hypotheses for sources of developmental difficulty in this task, namely, lack of quantifier knowledge and developing inhibitory control (Horowitz, Schneider, & Frank, in prep.). We reasoned that the Alternatives Hypothesis necessitates familiarity with and ability to contrast alternatives on the quantifier scale $\langle \textit{none} - \textit{some} - \textit{all} \rangle$; if children's quantifier knowledge is absent or unestablished, it might lead to failures in making an implicature. Another possible cause of children's struggles with scalar implicatures might be that children have complete quantifier knowledge, but are unable to inhibit an impulse to choose a more salient alternative (e.g., choosing the book with the most cats upon hearing "On the cover of

my book, *some* of the pictures are cats.”)

We explored these two hypotheses in an individual differences task, running our scalar implicature (Horowitz & Frank, 2015) task in conjunction with quantifier-knowledge (D. Barner, Chow, & Yang, 2009) and inhibitory control (Zelazo, 2006) tasks. Overall, we found that children’s scalar implicature performance was strongly correlated with quantifier knowledge, even when controlling for age. While inhibitory control was strongly correlated with age, it was not related to children’s performance on the SI task (Horowitz et al., in prep.).

Although we found a significant relationship between quantifier knowledge and performance on our scalar implicature task, it is likely that children’s difficulties in this task has more than one source. Importantly, in the course of conducting this individual differences study, we observed that children who succeeded both in making a scalar implicature and comprehending “none” seemed to have increased response latencies. This increase in reaction time may be an important indicator of how children are using and evaluating the salient alternatives.

It is possible that children who have a fully-established quantifier scale must take additional time in using this information to process a scalar implicature. This hypothesis has some support in previous literature; Huang & Snedeker (2009) explored online measures scalar implicature processing through eye-tracking, and found a delay in children’s location of the referent of a scalar implicature. The exact relationship between reaction time and accuracy using a supportive behavioral paradigm, however, is largely unknown. Here, we explore children’s behavioral response latencies in an iPad adaptation of Horowitz and Frank’s (2015) scalar implicature task. In this study, we explore our hypothesis that computing a scalar implicature might incur additional processing time for children as they contrast the relevant lexical alternatives to make a correct decision.

In our analyses, we explore overall accuracy and patterns of performance, as in (Horowitz & Frank, 2015; Horowitz et al., in prep.), and find that children not only struggle in making a scalar implicature, but also grapple with the quantifier “none” until fairly late in development. In examining reaction time patterns across all quantifier types, we find evidence of a speed-accuracy tradeoff associated with these two quantifiers, even later in development. Finally, we use a Drift Diffusion




Figure 1: Example trial stimuli used in Horowitz and Frank (2015).

Age group	N	Mean	Median	SD
3–3.5 years	22	3.27	3.27	0.15
3.5–4 years	34	3.78	3.73	0.15
4–4.5 years	24	4.29	4.29	0.14
4.5–5 years	30	4.76	4.76	0.15
5–6.5 years	24	5.55	5.56	0.36

Table 1: Age demographic information for all participants.

Model to our data to explore the source of this increased reaction time. Overall, our findings indicate that while quantifier knowledge is a key factor in successfully computing scalar implicatures, using this information to successfully compute a scalar implicature is particularly difficult, and seems to require additional processing.

Method

In this study, we adapted a scalar implicature paradigm developed by Horowitz and Frank (2015) for the iPad.¹ In addition to capturing detailed reaction time data, this version included more trials, and standardized prosody across all trials, in addition to a completely randomized design.²

Participants

Table 1 shows the breakdown of age information for all participants. Included in analyses are 134, 134, 134, 134, 134 children out of a planned sample of 120 participants. We ran 20 additional children, who were excluded from analysis for low English language exposure or $< \$50\%$ of trials completed. Included in our sample were 75 females and 59 males. Based on (Horowitz & Frank, 2015; Horowitz et al., in prep.), the initially planned sample size was 96 children from 3–5 years. After collecting data from 57 participants, however, we observed significantly lower performance on implicature trials across all age groups, indicating that the iPad adaptation of the scalar implicature task was slightly more challenging for all children, and included an older age group of 24 5–6.5-year-olds.

Stimuli

The general format of the task was identical to (Horowitz & Frank, 2015), with the exception of added items for additional trials. The study was programmed in HTML, CSS, and JavaScript, and displayed to children on a full-sized iPad. Each trial displayed three book covers, each containing a set of four familiar objects (Figure 1). Each trial allowed 2.5s for children to visually inspect the three book covers, before the experiment played the trial prompt (e.g., “On the cover of my book, *none* of the pictures are cats.”). Each trial was randomized, with the exception that similar items were displayed together (e.g., food, clothing). Each session involved

¹The full experiment can be viewed online at https://rosemshneider.github.io/tablet_exp/si-tablet.html.

²All of our data, processing, experimental stimuli, and analysis code can be viewed in the version control repository for this paper at: <https://github.com/rosemshneider/SI-tablet>.




Figure 2: Children’s overall accuracy for each quantifier type. Bars show mean performance for each age group. Error bars are 95 percent confidence intervals computed by non-parametric bootstrap.

30 trials, with 10 trials per quantifier-type (“all”, “some”, and “none”). In our randomization, quantifier triad order, items (within category), target item, and quantifier, were randomized for all participants.

Procedure

Sessions took place individually in a small testing room away from either the museum or the classroom. To familiarize children to the iPad, each session began with the “dot game,” which required them to press dots on the screen as fast as possible.

After the dot game, the experimenter introduced them to “Hannah,” a cartoon character who wanted to play a guessing game with her books. The experimenter explained that Hannah would show the child three books, and would give *one* hint about which book she had in mind. The experimenter emphasized that Hannah would only give one hint, so they had to listen carefully. Children then saw a practice trial with three books featuring a refrigerator, a TV, and a couch. After 2.5s, a female voice said “On the cover of my book there’s a TV.” Once children correctly made their selection, a green box appeared around the selection. Children moved trials along at their own pace by pressing a green button that appeared after they had made their selection.

Reaction times were measured from the onset of the target word. Each audio clip used the same three frames (e.g., “On the cover of my book, *some* of the pictures...”) so that prosody was emphasized equally across all trials. The average length of each audio clip (including target item phrase, e.g., “...are cats”) was approximately 6s. In all, there were 270 different target items and audio clips. Children could only make one selection. If a child was not paying attention, or if she did not hear Hannah’s prompt, the experimenter repeated it, matching the original prosody.

Results

In analyzing the results, we excluded any trials in which reaction time exceeded fifteen seconds, which indicated that the




Figure 3: Frequency histogram of correct responses for each trial type, across all participants.

child had missed the prompt, or was not paying attention. After this initial cut, we excluded responses outside three standard deviations of the log of mean reaction time. This cleaning process resulted in a data loss of 80 trials (2.07%). Our planned analyses for this study included explorations of accuracy, reaction time, and fitting a drift diffusion model (Ratcliff & Rouder, 1998) to our reaction time data.

Accuracy

In our first planned analysis, we explored children’s overall patterns of accuracy for each quantifier type. Figure 2 shows children’s for each trial type, split by age group. For each age group, we saw significantly lower accuracy for the quantifiers “some” and “none” in comparison to “all” in independent t-tests within each age group ($p < .01$ for all tests). These results replicate our previous findings using this paradigm (Horowitz & Frank, 2015; Horowitz et al., in prep.), indicating that our iPad version was a successful and appropriate adaptation of the scalar implicature task.

One difference from the previous results was in implicature trials. We found that children aged 3–5 years performed significantly lower on “some” (implicature) trials in this task in comparison with (Horowitz et al., in prep.) in independent t-tests ($p < .009$ for all tests). This difference is most likely due to the fact that our adaptation relied strictly on verbal communication, rather than other social and nonverbal cues.

Statistical modeling In exploring children’s significantly lower performance on “some” and “none” trials, we ran a logistic mixed effects model predicting correct response as an interaction of age and trial type, with random effects of trial type and participant.³ We found that performance was significantly lower on “some” ($\beta = -7.61$, $p < .0001$) and “none” trials ($\beta = -10.09$, $p < .0001$). We also found a significant interaction between age and trial type on “none” trials ($\beta = 1.64$, $p < .0001$), indicating that children’s performance with this difficult quantifier increased with age.

³Mixed effects model fit in R using the lme4 package. The model specifications were as follows: correct ~ age * trial type + (trial type | subject id).

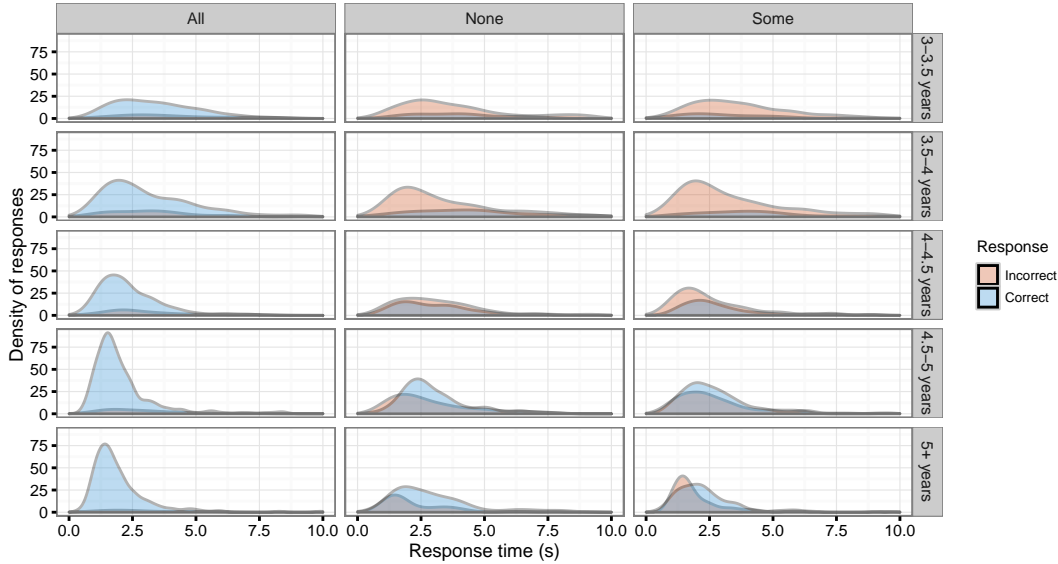


Figure 4: Density plots of reaction times for correct and incorrect responses on each trial type, split by age.

Correlation between “some” and “none” In previous research, a strong correlation has been found on children’s performance with the quantifiers “some” and “none” (Horowitz & Frank, 2015; Horowitz et al., in prep.). Figure 3 shows distributions of correct responses for all trial types. Once again, we found correlated performance between these two quantifiers ($r = 0.49$, $p < .0001$). In running Hartigan’s diptest for bimodality on these two quantifiers, we found significant bimodal distributions for “some” ($D = 0.08$, $p < .0001$) and “none” trials ($D = 0.11$, $p < .0001$). While the diptest also significantly rejected unimodality for “all” trials, this is most likely due to the distribution’s long tail. The observed significantly bimodal performance on “some” and “none” trials, however, indicate that while children’s accuracy is significantly worse for these quantifiers, they make responses in meaningful manner over the course of the study. In fact, although children’s performance in general was quite low on these trials, there were a number of children who got the majority of these trials correct (“Some”: $N = 28$; “None”: $N = 36$).

Discussion In our accuracy analyses, we found decreased performance across all age groups for the quantifiers “some” and “none,” and lower performance in general for “some” (implicature) trials in comparison to previous results (Horowitz & Frank, 2015; Horowitz et al., in prep.). Additionally, we also found bimodal and correlated performance in these two trial types. Finally, a statistical model revealed a main effects of age and trial type, as well as significant interactions between age and trial type, with children becoming more accurate on “some” and “none” trials with age.

Reaction time analyses

Response latencies are a largely-unexplored aspect of children’s ability to compute scalar implicatures. We hypothe-

sized that children’s reaction times on this task may reflect processing of a scalar prompt, and weighing quantifier alternatives. It is possible that this measure may provide a clue as to the nature of children’s correlated struggles with these terms. In recording reaction times, we began recording from the onset of the target nouns. Here, we explore overall trends in reaction times across this task, and the relationship between response latencies and accuracy on this task.

Developmental reaction time distribution Figure 4 shows the density of reaction times for each quantifier, faceted by age group and trial type. Overall, we found that reaction time was negatively correlated with age ($r = -0.25$, $p < .0001$). In exploring the relationship between accuracy and reaction time, we found preliminary evidence of a speed-accuracy tradeoff across all trial types. Figure ?? also shows evidence for increased response latencies associated with correct responses for the quantifiers “some” and “none”. Overall, we found that children were largely consistent in their performance across the course of the study in reaction time correlations between reaction times for “all” and “some” ($r = 0.76$, $p < .0001$), “all” and “none” ($r = 0.66$, $p < .0001$), and “some” and “none” trials ($r = 0.71$, $p < .0001$).

Statistical modeling We next turned to the relationship between age, reaction time, and accuracy. Our initial hypothesis was that successfully computing quantifiers might require additional processing time to contrast salient alternatives. In exploring this, we ran a planned linear mixed effects model predicting the log of reaction time as an interaction of age, trial number, and trial type with a random effect of trial type.⁴ We

⁴Mixed effects model fit in R using the lme4 package. The model specifications were as follows: $\log(\text{reaction time}) \sim \text{scale}(\text{age}) * \log(\text{trial number}) + \text{scale}(\text{age}) * \text{trial type} + (\text{trial type} | \text{subject id})$. We calculated p values by treating the t statistic as if it were a z statistic Barr, Levy,

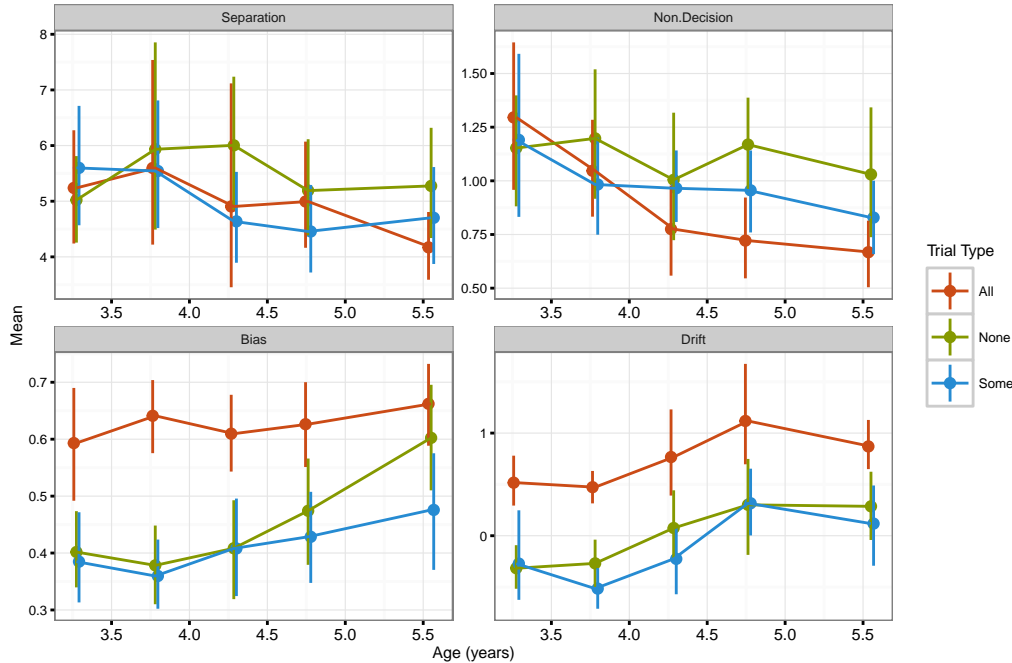


Figure 5: Parameter estimates for drift diffusion model, split by age and trial type. Error bars are 95 percent confidence intervals computed by nonparametric bootstrap.

found a main effect of trial number, with reaction times decreasing over the course of the study ($\beta = -0.1$, $p < .00001$), but found longer reaction times on “none” ($\beta = 0.23$, $p < .00001$), and “some” trials ($\beta = 0.1$, $p < .00001$). Interestingly, we also found an interaction between age and trial type, such that reaction times on “none” trials increased with age relative to “all” trials ($\beta = -0.02$, $p < .02$), and marginally (but not significantly) increased on “some” trials ($\beta = 0.14$, $p = .38$).

This interaction is particularly intriguing because in our previous accuracy model, we found increased performance on these trial types. While we find that older children are taking longer to respond to these trial types, they are more likely to answer correctly, even though reaction time is negatively correlated with age. This seems to indicate that there is a speed-accuracy tradeoff for these quantifiers.

Drift diffusion models

Our preliminary analysis of children’s reaction times in this task indicated greater response latencies associated with success on “some” and “none” trials. This suggests that children may be taking more time to actively compare and contrast scalar alternatives as they become more familiar with the quantifier scale. A drift diffusion model (DDM) can provide a more detailed view of the relationship between accuracy and reaction time in behavioral tasks (Milosavljevic, Malmaud, Huth, Koch, & Rangel, 2010; Ratcliff & Rouder, 1998).

Parameter estimation In DDM, a behavioral response (a correct or incorrect choice) is the result of noisy data accumulation through a diffusion process (operationalized by response time) (Ratcliff & Rouder, 1998). Responses have *separation boundaries* that are dependent on the amount of information needed to initiate a response, and *drift rate* formalizes the rate of data accumulation (Ratcliff & Rouder, 1998). An additional parameter of DDM is *nondecision*, which is the amount of time between encoding the stimuli, and initiating response. Finally, different responses may have a *bias*, or different starting point in the diffusion process, dependent on the stimuli (Ratcliff & Rouder, 1998).

In fitting a DDM to our data, we estimated parameters for each subject across all three trial types (“all”, “some”, and “none”) using the RWiener package. We then aggregated across subjects to obtain means and confidence intervals for each age group.⁵ Figure 5 shows the parameter estimates for each age group, split by trial type.

Using the DDM, we can explore in more depth the speed-accuracy tradeoff we observed in our statistical models. In our separation boundary estimates, we saw that over development the boundary for a correct response decreases, for all trial types, indicating that in general less data needs to be accumulated before making a decision in this task. This is also reflected in the bias estimates, which are very high for all age groups on “all” trials, but increase over development for “some” and “none” trials. Interestingly, we found that non-decision decreases markedly for “all” trials over devel-

Scheepers, & Tily (2013).

⁵Here we should address reaction time exclusion, when we’ve made a decision about it.

opment, but only slightly for “some” and “none” trials, while drift rates for all quantifiers increase over development.

In this model, we find that younger children spend less time processing a quantifier, and are also faster in making an incorrect decision with the difficult quantifiers “some” and “none”. Young children’s high drift rate and comparable nondecision estimates on “all” trials can be attributed to their tendency to select the alternative for “all” on the majority of trials. Later in development, however, we find that older children seem to spend more time processing the quantifiers “some” and “none”, and consequently are more likely to make a correct response. Intriguingly, we observe this trend even though older children have a greater bias and a lower separation boundary with these quantifier types.

General Discussion

Our primary question in this study centered on whether children who succeed on a scalar implicature task (and also understand the quantifiers involved) take more time to process. We adapted a previously validated scalar implicature task (Horowitz & Frank, 2015; Horowitz et al., in prep.) for the iPad to explore whether children’s success how response latencies interacted with accuracy in this task.

In our analyses, we replicated previous results (Horowitz & Frank, 2015; Horowitz et al., in prep.) in our finding that children were overall less accurate when evaluating the quantifiers “some” and “none” in comparison to “all,” but that their performance increased over development. We again found evidence of bimodal and correlated performance on these two quantifier types, suggesting a common source of difficulty.

In our extension of this paradigm, we collected reaction time data for these quantifier types to investigate the relationship between reaction time and accuracy. In our reaction time analyses, we found evidence of a speed-accuracy tradeoff, as well as an interaction between reaction time and age, with older children taking a slightly longer time to respond to these trials, but ultimately being more accurate.

Using a Drift Diffusion Model, we explored reaction time and accuracy patterns in more depth. Overall, the model predicted that children spend more time processing these difficult pragmatic implicatures as they get older, and are more likely to make a correct response as a result. Thus, it appears that success in making a scalar implicature is not only the result of being familiar with the quantifier scale, but also taking additional time to process and integrate in a pragmatic framework in order to make a response.

Our work contributes to the existing literature in utilizing a novel method to collect accurate and detailed reaction time data on a scalar implicature task. Response latencies are an important indicator of the pragmatic challenges that children face in processing implicatures. Additionally, our findings replicate previous work, providing evidence for the appropriateness of this paradigm in targeting scalar implicatures. Further, our larger sample size, increased number of trials, and randomized design strengthen our analytical power, and

allow for more detailed inferences from the data.

(Final paragraph on quantifier knowledge, inhibitory control, scalar implicatures, and how they relate more broadly to language).

Acknowledgements

Special thanks to Bing Nursery School, the San Jose Children’s Discovery Museum, Veronica Cristiano, Rachel Walker, and Tamara Mekler for their help with data collection.

References

- Barner, D., & Bachrach, A. (2010). Inference and exact numerical representation in early language development. *Cognitive Psychology*, 60(1), 40–62.
- Barner, D., Brooks, N., & Bale, A. (2011). Accessing the unsaid: The role of scalar alternatives in children’s pragmatic inference. *Cognition*, 118(1), 84–93.
- Barner, D., Chow, K., & Yang, S. (2009). Finding one’s meaning: A test of the relation between quantifiers and integers in language development. *Cognitive Psychology*, 58(2), 195–219.
- Barr, D. J., Levy, R., Scheepers, C., & Tily, H. (2013). Random effects structure for confirmatory hypothesis testing: Keep it maximal. *Journal of Memory and Language*, 68(3), 255–278.
- Grice, H. (1975). *Logic and conversation* (pp. 41–58).
- Guasti, M., Chierchia, G., Crain, S., Foppolo, F., Gualmini, A., & Meroni, L. (2005). Why children and adults sometimes (but not always) compute implicatures. *Language and Cognitive Processes*, 20(5), 667–696.
- Horowitz, A., & Frank, M. C. (2015). Sources of developmental change in pragmatic inferences about scalar terms. In *Proceedings of the 37th annual conference of the cognitive science society*.
- Horowitz, A., Schneider, R. M., & Frank, M. C. (in prep.). The trouble with quantifiers: Children’s difficulties with “some” and “none”.
- Huang, Y. T., & Snedeker, J. (2009). Online interpretation of scalar quantifiers: Insight into the semantics-pragmatics interface. *Cognitive Psychology*, 58(3), 376–415.
- Katsos, N., & Bishop, D. (2011). Pragmatic tolerance: Implications for the acquisition of informativeness and implicature. *Cognition*, 120(1), 67–81.
- Milosavljevic, M., Malmaud, J., Huth, A., Koch, A., & Rangel, A. (2010). Drift diffusion model can account for accuracy and reaction time of value-based choices under high and low time pressure. *Judgment and Decision Making*, 5(6), 437–449.
- Noveck, I. (2001). When children are more logical than adults: Experimental investigations of scalar implicature. *Cognition*, 78(2), 165–188.
- Papafragou, A., & Musolino, J. (2003). Scalar implicatures: Experiments at the semantics-pragmatics interface. *Cognition*, 86(3), 253–282.

- Papafragou, A., & Tantalou, N. (2004). Children's computation of implicatures. *Language Acquisition*, 12, 71–82.
- Ratcliff, R., & Rouder, J. (1998). Modeling response times for two-choice decisions. *Psychological Science*, 9(5), 347–356.
- Stiller, A., Goodman, N., & Frank, M. (2014). Ad-hoc implicature in preschool children. *Language Learning and Development*, 11(2), 176–190.
- Zelazo, P. D. (2006). The dimensional change card sort (dCCS): A method of assessing executive function in children. *Nature Protocols*, 1, 297–301.